

LLVM Fence Synthesis

PARALLEL AND CONCURRENT PROGRAMMING

CS4560

Authors:

Group 3

Mihnea Bernevig (5496004)
Yigit Colakoglu (5495121)

April 13, 2025

Contents

1	Overview	2
2	Fence Insertion	2
2.1	TSO vs. PSO Fence Synthesis	2
3	Fence Optimization	5
4	Testing	6

1 Overview

In this project, we implemented three transformation passes for the LLVM compiler. Two of these passes are used to greedily insert Sequentially-Consistent (SC) fences for different target memory models, while the third pass eliminates redundant fences to optimize for overall code size.

Sequentially-Consistent fences enforce a happens-before relationship between the operations preceding the fence and those following it. This enforcement guarantees that all operations before the fence are completed prior to the execution of those after it, both at the hardware and compiler levels. We insert the fences during the IR stage so the LLVM backend can convert the IR fence instructions to the target architecture, or completely ignore them if necessary (For instance, SC fences are essentially ignored when compiling to x86).

2 Fence Insertion

We have developed two separate passes for fence insertion, one ensuring consistency under the Total Store Order (TSO) model and the other under the Partial Store Order (PSO) model.

Both of these passes work by traversing each function starting from their initial basic block, traversing all possible control flows of a program while keeping track of the last atomic memory operation that was performed so that it can insert a fence if a consecutive operation pair that is "too relaxed" appears in the path. The method used to traverse the basic blocks and insert fences for TSO and PSO memory order is shown in Algorithms 1 and 2, respectively.

Both of our algorithms insert two fences per operation pair, once after the initial operation, and once before the second operation. We have found that this approach works better in practice, especially when the LLVM IR contains ϕ nodes.

2.1 TSO vs. PSO Fence Synthesis

The main difference between the TSO and PSO memory synthesis passes is the consecutive atomic memory operation pairs that are "allowed". For instance, the PSO memory model allows ww memory operations to be re-ordered while TSO does not. So, the TSO pass inserts a fence before the second write if it encounters such a pair with relaxed memory order while the PSO pass does not care. A more detailed enumeration of the transformations we perform for each target memory order can be found in Table 1.

Op ₁	Order ₁	Op ₂	Order ₂	TSO	PSO
R	RLX	R	RLX	R-Fence-R	R-Fence-R
R	RLX	W	RLX	R-Fence-W	R-Fence-W
W	RLX	R	RLX	R-W	R-W
W	RLX	W	RLX	W-Fence-W	W-W
R	ACQ	*	*	R-Op ₂	R-Op ₂
*	*	W	REL	Op ₁ -W	Op ₁ -W
W	REL	R	ACQ	W-R	W-R
*	SEQ_CST	*	SEQ_CST	Op ₁ -Op ₂	Op ₁ -Op ₂
*	ACQ_REL	*	ACQ_REL	Op ₁ -Op ₂	Op ₁ -Op ₂

Table 1: Memory operation ordering under TSO and PSO

TL;DR: Fences are only introduced between two relaxed atomic memory operations as the C11 memory ordering for operations with other memory orders already behave the way we want them to.

Algorithm 1: TraverseBBGraph

Input: Basic block BB , ordering $order$, last memory operation $lastMemOp$

```
1 if  $BB$  is empty then
2   return
3 for each instruction  $I$  in  $BB$  do
4   if  $I$  is a  $LoadInst$  then
5      $loadOrder \leftarrow I.getOrdering()$ ;
6     if  $lastMemOp$  is null then
7        $lastMemOp \leftarrow I$ ;
8       if  $order = Unordered \wedge loadOrder \in \{Unordered, Monotonic\}$  then
9         insert fences with Sequentially-Consistent ordering before  $I$ ;
10      continue;
11    if  $loadOrder \notin \{Unordered, Monotonic, Acquire\}$  // Skip unsupported ordering then
12       $order \leftarrow Unordered$ ;
13       $lastMemOp \leftarrow I$ ;
14      continue;
15    if  $lastMemOp$  is a  $StoreInst$  then
16      // No fence needed when previous operation is a store and current is a load (Write-Read)
17    if  $lastMemOp$  is a  $LoadInst$  and  $order \notin \{Acquire, AcquireRelease, SequentiallyConsistent\}$ 
18      then
19        insert fences with Sequentially-Consistent ordering before  $I$  and after  $lastMemOp$ ;
20       $order \leftarrow loadOrder$ ;
21       $lastMemOp \leftarrow I$ ;
22    else if  $I$  is a  $StoreInst$  then
23       $storeOrder \leftarrow I.getOrdering()$ ;
24      if  $lastMemOp$  is null then
25         $lastMemOp \leftarrow I$ ;
26        if  $order = Unordered \wedge storeOrder \in \{Unordered, Monotonic\}$  then
27          insert fences with Sequentially-Consistent ordering before  $I$ ;
28          continue;
29        if  $lastMemOp$  is a  $StoreInst$  and  $order \notin \{Release, AcquireRelease, SequentiallyConsistent\}$ 
30          then
31            insert fences with Sequentially-Consistent ordering before  $I$  and after  $lastMemOp$ ;
32          else if  $lastMemOp$  is a  $LoadInst$  and
33             $order \notin \{Acquire, AcquireRelease, SequentiallyConsistent\}$  then
34              insert fences with Sequentially-Consistent ordering before  $I$  and after  $lastMemOp$ ;
35             $order \leftarrow storeOrder$ ;
36             $lastMemOp \leftarrow I$ ;
37    else if  $I$  is a  $FenceInst$  then
38       $order \leftarrow I.getOrdering()$ ;
39    else if  $I$  is a  $Terminator$  instruction then
40      if  $I$  is a  $BranchInst$  then
41        for each successor  $BB'$  of  $I$  do
42          if  $BB' = BB$  then
43            continue;
44           $TraverseBBGraph(BB', order, lastMemOp)$ ;
45      else if  $I$  is a  $SwitchInst$  then
46        for each successor  $BB'$  of  $I$  do
47          if  $BB' = BB$  then
48            continue;
49           $TraverseBBGraph(BB', order, lastMemOp)$ ;
50    else
51      return;
```

Algorithm 2: TraverseBBGraphPSO

Input: BasicBlock BB , AtomicOrdering $order$, Instruction pointer $lastMemOp$

```
1 if  $BB$  is empty then
2   return
3 for each instruction  $I$  in  $BB$  do
4   if  $I$  is a LoadInst then
5      $loadOrder \leftarrow I.getOrdering()$ ;
6     if  $lastMemOp$  is null then
7        $lastMemOp \leftarrow I$ ;
8       if  $order = \text{Unordered} \wedge loadOrder \in \{\text{Unordered}, \text{Monotonic}\}$  then
9         insert fences with Sequentially-Consistent ordering before  $I$ ;
10      continue;
11    if  $loadOrder \notin \{\text{Unordered}, \text{Monotonic}, \text{Acquire}\}$  then
12       $order \leftarrow \text{Unordered}$ ;
13       $lastMemOp \leftarrow I$ ;
14      continue;
15    if  $lastMemOp$  is a StoreInst then
16      continue;
17    if  $lastMemOp$  is a LoadInst and  $order \notin \{\text{Acquire}, \text{AcquireRelease}, \text{SequentiallyConsistent}\}$ 
18      then
19        insert fences with Sequentially-Consistent ordering before  $I$  and after  $lastMemOp$ ;
20         $order \leftarrow loadOrder$ ;
21         $lastMemOp \leftarrow I$ ;
22    else if  $I$  is a StoreInst then
23       $storeOrder \leftarrow I.getOrdering()$ ;
24      if  $lastMemOp$  is null then
25         $lastMemOp \leftarrow I$ ;
26        if  $order = \text{Unordered} \wedge storeOrder \in \{\text{Unordered}, \text{Monotonic}\}$  then
27          insert fences with Sequentially-Consistent ordering before  $I$ ;
28          continue;
29        if  $lastMemOp$  is a StoreInst then
30          insert fences with Sequentially-Consistent ordering before  $I$  and after  $lastMemOp$ ;
31        else if  $lastMemOp$  is a LoadInst and
32           $order \notin \{\text{Acquire}, \text{AcquireRelease}, \text{SequentiallyConsistent}\}$  then
33            insert fences with Sequentially-Consistent ordering before  $I$  and after  $lastMemOp$ ;
34             $order \leftarrow storeOrder$ ;
35             $lastMemOp \leftarrow I$ ;
36    else if  $I$  is a FenceInst then
37       $order \leftarrow I.getOrdering()$ ;
38    else if  $I$  is a Terminator instruction then
39      if  $I$  is a BranchInst then
40        for each successor  $BB'$  of  $I$  do
41          if  $BB' = BB$  then
42            continue;
43           $TraverseBBGraphPSO(BB', order, lastMemOp)$ ;
44      else if  $I$  is a SwitchInst then
45        for each successor  $BB'$  of  $I$  do
46          if  $BB' = BB$  then
47            continue;
48           $TraverseBBGraphPSO(BB', order, lastMemOp)$ ;
49    else
50      return;
```

3 Fence Optimization

As we mentioned earlier, our initial methods insert fences greedily without attempting any optimizations to not insert any redundant fences. After this synthesis stage, we implemented a secondary optimization to remove redundant fences using redundant fence elimination techniques using the min-cut of the control flow graph [1].

In order to remove redundant fences, we build a special graph by finding and removing each fence instruction in the control flow, and inserting a node before and after where they used to be. Afterwards, for each inserted node before/after the fence, we traverse the basic block graph upwards and downwards until we encounter a memory operation or the start/end of the function flow graph, respectively. Once one of the following nodes are inserted, they are connected to a special source (if before the fence) or sink (if after the fence) node, essentially forming a network flow graph for the function. The resulting graph is directed, with weights of 1 for all edges that are not connected to the source/sink and ∞ for the edges that are connected to the source/sink. The pseudocode for the graph generation functions can be found in Algorithms 3-5.

Once we have the network flow graph G for a function, we need to calculate the min-cut of the resulting graph, the edges of which will correspond to the optimal placement of graphs without changing the function's behavior. To calculate the min-cut, we run the Ford-Fulkerson max-flow algorithm on the graph G . This results in a residual graph G' with the edges on the min-cut having been fully saturated. This means (given G was connected as a result of our graph generation procedure), G' is split into two subgraphs $G_s = (V_s, E_s)$, $G_t = (V_t, E_t)$, with the edges of weight 0 between them corresponding to the min-cut. To find those edges, we traverse all nodes accessible from the source, marking all the vertices of the encounter. The edges whose out-vertex is marked and in-vertex is not are exactly the edges $MinCut = \{e = (a, b) | a \in V_s \wedge b \in V_t\}$. Finally, we insert back Sequential Consistency fences on the locations that correspond to the min-cut.

Algorithm 3: makeGraphUpwards

Input: Instruction pointer *root*, Graph *graph*

Output: Node pointer

```
1 bb  $\leftarrow$  root.getParent();
2 foundRoot  $\leftarrow$  false;
3 for each instruction inst in bb in reverse order do
4   if inst = root then
5     foundRoot  $\leftarrow$  true;
6     continue;
7   if  $\neg$ foundRoot then
8     continue;
9   if inst is a memory access (i.e., a LoadInst or StoreInst) then
10    node  $\leftarrow$  getNode(inst, after);
11    graph.addNode(node);
12    graph.addEdge(graph.source, node);
13    return node;
14 node  $\leftarrow$  getNodeAtBeginning(bb);
15 if bb = bb.getParent().getEntryBlock() then
16   graph.addNode(node);
17   graph.addEdge(graph.source, node);
18   return node;
19 for each predecessor pred of bb do
20   node2  $\leftarrow$  getNodeAtEnd(pred);
21   graph.addNode(node2);
22   lastInst  $\leftarrow$  getLastInst(pred);
23   node3  $\leftarrow$  makeGraphUpwards(lastInst, graph);
24   if node3  $\neq$  null then
25     graph.addEdge(node, node2);
26     graph.addNode(node3);
27     graph.addEdge(node2, node3);
28 return node;
```

Algorithm 4: makeGraphDownwards

Input: Instruction pointer *root*, Graph *graph*

Output: Node pointer

```
1 bb ← root.getParent();
2 foundRoot ← false;
3 for each instruction inst in bb (in order) do
4   if inst = root then
5     | foundRoot ← true;
6     | continue;
7   if ¬foundRoot then
8     | continue;
9   if inst is a memory access (i.e., a LoadInst, StoreInst, or ReturnInst) then
10    | node ← getNode(inst, before);
11    | graph.addNode(node);
12    | graph.addEdge(node, graph.sink);
13    | return node;
14 node ← getNodeAtEnd(bb);
15 for each successor succ of bb do
16   node2 ← getNodeAtBeginning(succ);
17   graph.addNode(node2);
18   firstInst ← getFirstInst(succ);
19   node3 ← makeGraphDownwards(firstInst, graph);
20   if node3 ≠ null then
21     | graph.addEdge(node, node2);
22     | graph.addNode(node3);
23     | graph.addEdge(node2, node3);
24 return node;
```

Algorithm 5: TransformFunction

Input: Function pointer *fun*, Graph *graph*

```
1 for each Basic Block bb in fun do
2   for each Instruction inst in bb do
3     if inst is a FenceInst then
4       | nodeBeforeFence ← makeGraphUpwards(inst, graph);
5       | nodeAfterFence ← makeGraphDownwards(inst, graph);
6       | if nodeBeforeFence ≠ null and nodeAfterFence ≠ null then
7         | | graph.addNode(nodeBeforeFence);
8         | | graph.addNode(nodeAfterFence);
9         | | graph.addEdge(nodeBeforeFence, nodeAfterFence);
10 for each Basic Block bb in fun do
11   for each Instruction inst in bb do
12     if inst is a FenceInst then
13       | remove inst from bb;
```

4 Testing

In order to test the validity of our fence synthesis techniques we decided write litmus tests used in the lectures in LLVM IR and use the LLVM test suite in order to check that they were inserted in where we expected them to. We implemented three different litmus tests that contained consecutive read-write, write-write and write-read operations, which covered the insertion cases for both TSO and PSO's fence synthesis methods. More specifically, we implemented the load-buffer, store-buffer and message passing litmus tests.

For each litmus test, we tested the operations with different memory orders to check that the memory orders of operations are considered accordingly during the synthesis as well. Moreover, we implemented an additional test to ensure that the control flow is accounted for when inserting fences.

References

- [1] R. Morisset and F. Zappa Nardelli, “Partially redundant fence elimination for x86, arm, and power processors,” in *Proceedings of the 26th International Conference on Compiler Construction*, ser. CC 2017, Austin, TX, USA: Association for Computing Machinery, 2017, pp. 1–10, ISBN: 9781450352338. DOI: [10.1145/3033019.3033021](https://doi.org/10.1145/3033019.3033021). [Online]. Available: <https://doi.org/10.1145/3033019.3033021>.